

CS-GY 6233 Operating systems - Final Project

📅 Date	@December 8, 2021
☰ Team mate 1	Rigveda Vangipurapu rv2205
☰ Team mate 2	Jaya Amit Sai Gurralla jg6660

Basics

We have written a program using system calls - open, read, write and close.

The program reads the `block_size` number of characters each time, and reads `block_count` number of blocks.

So the total `file_size` being read by our code is `block_size * block_count`

Our program returns the XOR value of whatever is read so as to make sure the read operation is taking place correctly.

Usage: `./run <filename> [-r|-w] <block_size> <block_count>`

Output:

```
XOR value: 0000000e
rigvedavangipurapu@Rigvedas-Air OSfinal % gcc run.c -o run
rigvedavangipurapu@Rigvedas-Air OSfinal % ./run dumb.txt -r 1000 10
XOR value: 0000000e
```



Take away from this task: We learnt the usage of these system calls and figured out boiler-plate stuff like how to declare buffers etc

Measurement

We wrote a program to measure the performance of the program.

Our implementation to the same was as follows:

1. Open the file given as input in the command line.
2. Set timer and block_count to 0.
3. Start reading block_size amount of information from it.
4. Each time one block is read, the time taken is checked.
5. If the cumulative time taken does not exceed 5 seconds, continue the read operation and increment the block_count.
6. If the cumulative time taken reaches (or exceeds) 5 seconds, stop reading.
7. Report the file size to be block_count*block size

Usage: ./run2 <filename> <block_size>

Output:

```
(base) rigvedavangipurapu@Rigvedas-Air OSfinal % gcc run2.c -o run2
(base) rigvedavangipurapu@Rigvedas-Air OSfinal % ./run2 dumb.txt 1024

File size: 32768.000000 MB
```

Extra credit task:

We wrote a program which reads data from a file as blocks, and writes it to another file.

Usage of dd program: dd if=<input_file> of=<output_file>

Usage of our program: ./compareddd <input_file> <output_file> <block_size>

At block size 512, this is how the program compares with the dd program:

```

rigvedavangipurapu@Rigvedas-Air OSfinal % ./comparedd dumb.txt hello.txt 512

file size 1048576 bytes
read: time taken 0.014670 secs
Performance at 512= 71478084.041996 bytes/s
rigvedavangipurapu@Rigvedas-Air OSfinal % dd if=dumb.txt of=hello.txt
2048+0 records in
2048+0 records out
1048576 bytes transferred in 0.011072 secs (94705883 bytes/sec)

```

*There is a minor variation each time the programs are run possibly due to caching and various other factors, this is one of the instances captured for reference.

Running the dd program and our program simultaneously, we noted the following observations:

File size	dd program: time taken	dd program: performance	our program: time taken	our program: performance
80 bytes	0.000052 secs	1539194 bytes/sec	0.000595 secs	3442859.555912 bytes/s
1MB	0.010605 secs	98874722 bytes/sec	0.009082 secs	115455503.927336 bytes/sec
1GB	17.181206 secs	164059390 bytes/sec	18.279118 secs	154205370.663034 bytes/sec

Raw Performance

This program is yet another way to measure the performance of the program.

We witness how the block size affects the performance.

We read the [ubuntu-21.04-desktop-amd64.iso](#) file which is sufficiently large to accomodate the block sizes we use.

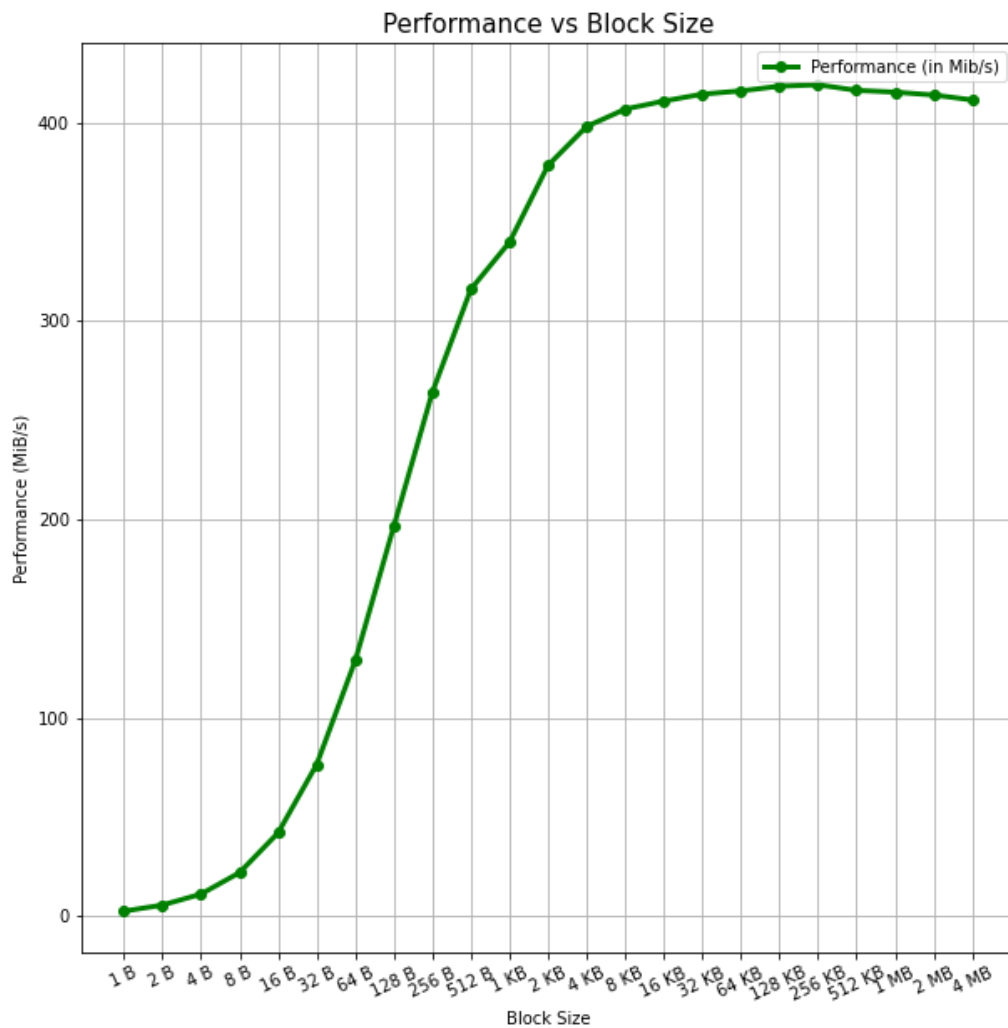
Now we changed the block sizes starting from 1 Byte ranging till 4 MB.

The following values were obtained:

Block size	Performance (MiB/s)
1 B	2.675

Block size	Performance (MiB/s)
2 B	5.754
4 B	11.30
8 B	22.110
16 B	42.153
32 B	76.597
64 B	128.894
128 B	196.524
256 B	264.083
512 B	315.987
1 KB	339.620
2 KB	378.327
4 KB	397.877
8 KB	406.633
16 KB	410.70
32 KB	414.189
64 KB	415.818
128 KB	418.186
256 KB	418.896
512 KB	416.103
1 MB	415.190
2 MB	413.76
4 MB	411.279

Plotting the above values, we obtain the following graph:



Usage: ./performancegraph <block_size>

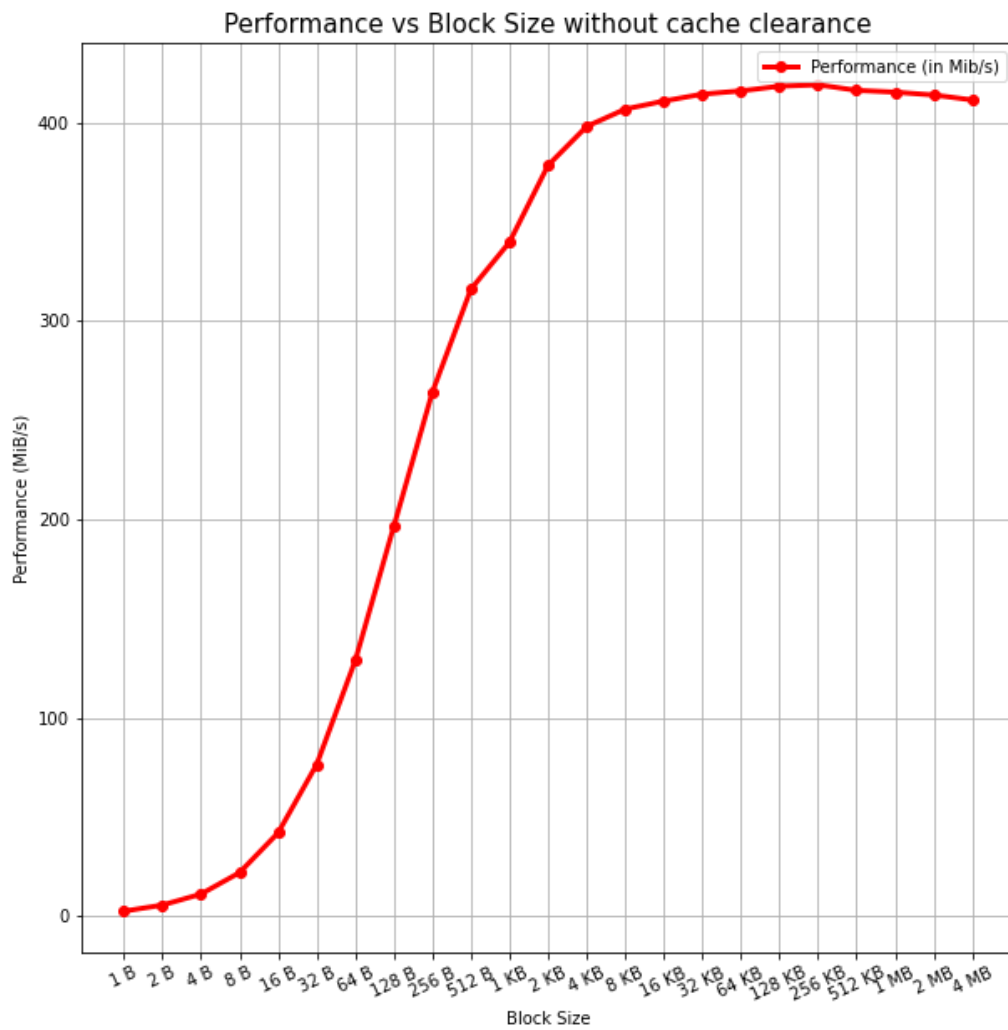
Caching

We ran the performance measurement code using the same [ubuntu-21.04-desktop-amd64.iso](#) file with and without clearing the cache to see the results.

We observed the following results:

Without clearing cache:


Block size	Performance (MiB/s)
1 B	2.675
2 B	5.754
4 B	11.30
8 B	22.110
16 B	42.153
32 B	76.597
64 B	128.894
128 B	196.524
256 B	264.083
512 B	315.987
1 KB	339.620
2 KB	378.327
4 KB	397.877
8 KB	406.633
16 KB	410.70
32 KB	414.189
64 KB	415.818
128 KB	418.186
256 KB	418.896
512 KB	416.103
1 MB	415.190
2 MB	413.76
4 MB	411.279



With clearing cache each time:

Block size	Performance (MiB/s)
1 B	2.023
2 B	3.456
4 B	6.567
8 B	11.309
16 B	21.821
32 B	39.532
64 B	61.123

Block size	Performance (MiB/s)
128 B	81.793
256 B	87.440
512 B	89.413
1 KB	88.562
2 KB	88.434
4 KB	87.464
8 KB	87.568
16 KB	91.200
32 KB	83.375
64 KB	87.633
128 KB	88.002
256 KB	88.662
512 KB	88.901
1 MB	89.012
2 MB	89.102
4 MB	89.110



Take away: When the program is executed by clearing the cache, every time the data is fetched from the RAM, resulting in larger read times and lesser performance.

On the other hand, if the data is cached, it is much easier for the CPU to fetch the data and carryout the execution, resulting in smaller read times and greater performance.

Extra Credit

```
sudo sh -c "/usr/bin/echo 3 > /proc/sys/vm/drop_caches"
```

There are 3 options to clear the cache without interrupting any processes or services:

1. echo 1 : to clear page cache only
2. echo 2 : to clear dentries and inodes
3. echo 3 : to clear page cache, dentries and inodes

An **Inode** is a data structure which provides a representation of a file.

Dentries is a data structure which represents a directory or a folder. So, dentries can be used to store cache, if they exist then store or check from memory.

The **Page cache** is anything the OS could store in the memory or hold it rather read it from a file.

System Calls

Keeping block size as 1 Byte, we tried various system calls to see how many system calls can be performed in a second.

We experimented with the following system calls:

read

write

open

lseek

System call	File Size	Time taken (secs)	Performance (B/s) (Number of system calls per second)	Performance (MiB/s)
read	1MB	0.431471	2430234.556102 B/s	2.317652 MiB/s
write	1MB	1.762226	595029.206060 B/s	0.567464 MiB/s
open	1MB	0.267061	3926353.976176 B/s	3.744463 MiB/s
lseek	1MB	0.188419	5565132.763460 B/s	5.307324 MiB/s

Output:

```
rigvedavangipurapu@Rigvedas-Air OSfinal % ./systemcalls hello.txt 1
```

```
file size 1048576
read: time taken 0.431471
read: performance at 1= 2.317652 MiB/s
read: performance at 1= 2430234.556102 B/s
```

```
file size 1048576
lseek: time taken 0.188419
lseek: performance at 1= 5.307324 MiB/s
lseek: performance at 1= 5565132.763460 B/s
```

```
file size 1048576
write: time taken 1.762226
write: performance at 1= 0.567464 MiB/s
write: performance at 1= 595029.206060 B/s
```

```
file size 1048576
open: time taken 0.267061
open: performance at 1= 3.744463 MiB/s
open: performance at 1= 3926353.976176 B/s%
```



Take away from this task: We observed that the number of system calls needed for block size=1 is very high, which means a very large number of I/O operations.

Since I/O operations are costly, the time taken to read a file is very large given block size = 1

Raw Performance

Based on the previous observations, we chose our block size to be: **256KB**

Usage: ./fast <file name>

Output:

```
(base) rigvedavangipurapu@Rigvedas-Air OSfinal % gcc fast.c -o fast
(base) rigvedavangipurapu@Rigvedas-Air OSfinal % ./fast ubuntu-21.04-desktop-amd64.iso
XOR value is a7eeb2d9
Time taken to run 262144 is 6.545756

Block count = 10753
Speed = 410.685955 MiB/s
(base) rigvedavangipurapu@Rigvedas-Air OSfinal % █
```

Expected performance:

Block size: 256 KB	Performance (MiB/s)
Without clearing cache	418.896
After clearing cache	88.662